

HW 4

March 4, 2021

1 IST 387 HW 4

Copyright 2021, Jeffrey Stanton, Jeffrey Saltz, and Jasmina Tacheva

```
[3]: # Enter your name here: Connor Hanan
```

1.0.1 Attribution statement: (choose only one and delete the rest)

```
[4]: # 1. I did this homework by myself, with help from the book and the professor.
```

(Chapters 8, 9, and 10 of Introduction to Data Science)

Reminders of things to practice from previous weeks: Descriptive statistics: `mean()` `max()` `min()` Sequence operator: `:` (For example, `1:4` is shorthand for `1, 2, 3, 4`) Create a function: `myFunc <- function(myArg) { } ?command: Ask R for help with a command`

This module: Sampling is a process of **drawing elements from a larger set**. In data science, when analysts work with data, they often work with a sample of the data, rather than all of the data (which we call the **population**), because of the expense of obtaining all of the data.

One must be careful, however, because **statistics from a sample rarely match the characteristics of the population**. The **goal of this homework** is to **sample from a data set several times and explore the meaning of the results**. Before you get started make sure to read Chapters 8-10 of An Introduction to Data Science. Don't forget your comments!

1.1 Part 1: Write a function to compute statistics for a vector of numeric values

- A. Create a new function which takes a numeric vector as its input argument and returns a list of statistics about that vector as the output. As a start, it should return the min, mean, and max of the vector. The function should be called **vectorStats**:

```
[18]: vectorStats <- function(n_vector) {  
  a <- min(n_vector)  
  b <- mean(n_vector)  
  c <- max(n_vector)  
  return (c(a,b,c))  
}
```

- B. Test your function by calling it with the numbers **one through ten**:

```
[19]: vectorStats(1:10)
```

```
1. 1 2. 5.5 3. 10
```

C. Enhance the `vectorStats()` function to add the **median** and **standard deviation** to the output.

```
[20]: vectorStats <- function(n_vector) {  
  a <- min(n_vector)  
  b <- mean(n_vector)  
  c <- max(n_vector)  
  d <- median(n_vector)  
  e <- sd(n_vector)  
  return (c(a,b,c,d,e))  
}
```

D. Retest your enhanced function by calling it with the numbers **one through ten**:

```
[21]: vectorStats(1:10)
```

```
1. 1 2. 5.5 3. 10 4. 5.5 5. 3.02765035409749
```

1.2 Part 2: Sample repeatedly from the mtcars built-in dataframe

E. Copy the mtcars dataframe:

```
[22]: myCars <- mtcars
```

Use `head(myCars)` and `tail(myCars)` to show the data. Add a comment that describes what each variable in the data set contains. **Hint:** Use the `?` or `help()` command with `mtcars` to get help on this dataset.

```
[26]: head(myCars)  
tail(myCars)  
  
#mpg is Miles/(US) gallon  
#cyl is Number of cylinders  
#disp is Displacement (cu.in.)  
#hp is Gross horsepower  
#drat is Rear axle ratio  
#wt is Weight (1000 lbs)  
#qsec is 1/4 mile time  
#vs is Engine (0 = V-shaped, 1 = straight)  
#am is Transmission (0 = automatic, 1 = manual)  
#gear is Number of forward gears  
#carb is Number of carburetors
```

		mpg <dbl>	cyl <dbl>	disp <dbl>	hp <dbl>	drat <dbl>	wt <dbl>	qsec <dbl>	vs <dbl>
A data.frame: 6 × 11	Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0
	Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0
	Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1
	Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1
	Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0
	Valiant	18.1	6	225	105	2.76	3.460	20.22	1

		mpg <dbl>	cyl <dbl>	disp <dbl>	hp <dbl>	drat <dbl>	wt <dbl>	qsec <dbl>	vs <dbl>
A data.frame: 6 × 11	Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.7	0
	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.9	1
	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.5	0
	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.5	0
	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.6	0
	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.6	1

F. Sample three observations from `myCars$mpg`.

```
[35]: sample(myCars$mpg, size=3, replace = T)
```

G. Call your `vectorStats()` function with a sample of three observations from `myCars$mpg`.

```
[37]: vectorStats(sample(myCars$mpg, size=3, replace = T))
```

```
1. 21 2. 24.26666666666667 3. 30.4 4. 21.4 5. 5.31538647074071
```

H. Use the `replicate()` function to repeat your sampling ten times. The first argument to `replicate()` is the number of repeats you want. The second argument is the little chunk of code you want repeated.

```
[73]: replicate(10, vectorStats(sample(myCars$mpg, size=3, replace = T)))
```

```

      10.40000  16.400000  10.400000  15.0000000  15.20000  10.400000  15.800000
      16.03333  21.733333  18.800000  15.7333333  21.43333  15.500000  21.800000
A matrix: 5 × 10 of type dbl 27.30000  24.400000  27.300000  16.4000000  33.90000  22.800000  30.400000
      10.40000  24.400000  18.700000  15.8000000  15.20000  13.300000  19.200000
      9.75722  4.618802  8.450444  0.7023769  10.79645  6.486139  7.639372
```

I. Write a comment describing why every replication produces a different result.

```
[42]: #every time it goes to replicate, it runs the sampling again, so it will be
      ↪ gathering a different set of 3 numbers from
      #myCars$mpg
```

J. Rerun your replication, this time doing 1000 replications and storing the output of `replicate()` in a variable called `values`.


```
[75]: values <- (replicate(1000, vectorStats(sample(myCars$mpg, size=3, replace = T))))
      values
```

	10.400000	14.30000	17.300000	21.400000	18.100000	10.400000	10.400000
	16.900000	18.50000	21.433333	24.400000	24.333333	16.366667	15.966667
A matrix: 5 × 1000 of type dbl	26.000000	26.00000	26.000000	30.400000	33.900000	24.400000	22.800000
	14.300000	15.20000	21.000000	21.400000	21.000000	14.300000	14.700000
	8.118497	6.51076	4.366158	5.196152	8.410906	7.225187	6.296295

K. Generate a **histogram** of the means stored in values. You need to **index into values** for that.

```
[76]: library(tidyverse)
```


```
[77]: ggplot(data.frame(values[2,]),aes(x=values[2,])) +  
  geom_histogram(bins=15,fill='black',col='grey')
```



L. Repeat the replicated sampling, but this time, raise your sample size from **3 to 22**. How does that affect your histogram? Explain in a comment.

```
[78]: values_expanded <- (replicate(1000, vectorStats(sample(myCars$mpg, size=22,
  ↳ replace = T))))
ggplot(data.frame(values_expanded[2,]), aes(x=values_expanded[2,])) +
  ↳ geom_histogram(bins=15, fill='black', col='grey')

#the histogram has moved towards a slightly more normal distribution
#increasing the sample size betters the accuracy of the mean, allowing fewer
  ↳ extremes to occur
```



```
[ ]:
```